

Quantum Machine Learning, Homework 2

- Author: Yan-Tong Lin, 0712238, National Chiao Tung University
- Lecturer: Dr. Alexey Melnikov, Principal Investigator, Terra Quantum
- Online Version at <https://hackmd.io/MkTt8-07SiWvv-p0LqR6Pw> (<https://hackmd.io/MkTt8-07SiWvv-p0LqR6Pw>).
- Jan 3rd, 2021

- Quantum Machine Learning, Homework 2
 - Note for Day 3 (12/4)
 - Overview
 - PS (cont.)
 - PS agent on Quantum System
 - Note for Day 4 (12/5)
 - Overview
 - PS code study
 - Note for Day 5 (12/11)
 - Overview
 - Quantum Speed up for PS
 - Quantum Communication Tasks with PS
 - Note for Day 6 (12/12)
 - Materials
 - HW2 Problem Set
 - Q1
 - A1
 - Q2
 - A2
 - Q3
 - A3

Note for Day 3 (12/4)

Overview

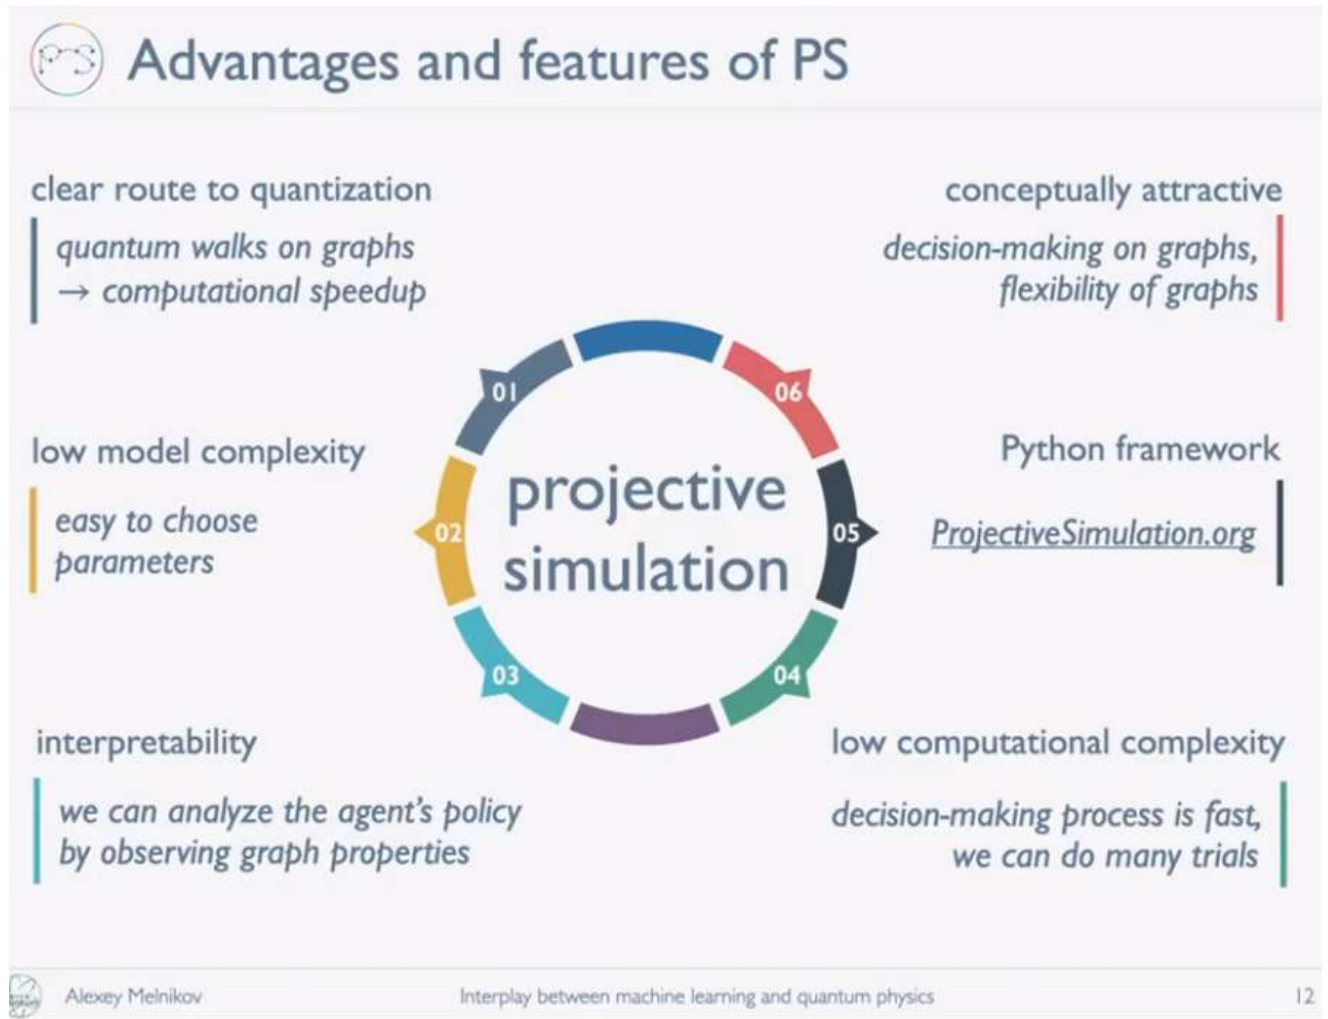
- More on PS

- analytic solution
- cost efficiency
- meta-param net
- applications
- PS applied in quantum labs
 - definitions
 - the task is hard
 - the model
 - the results

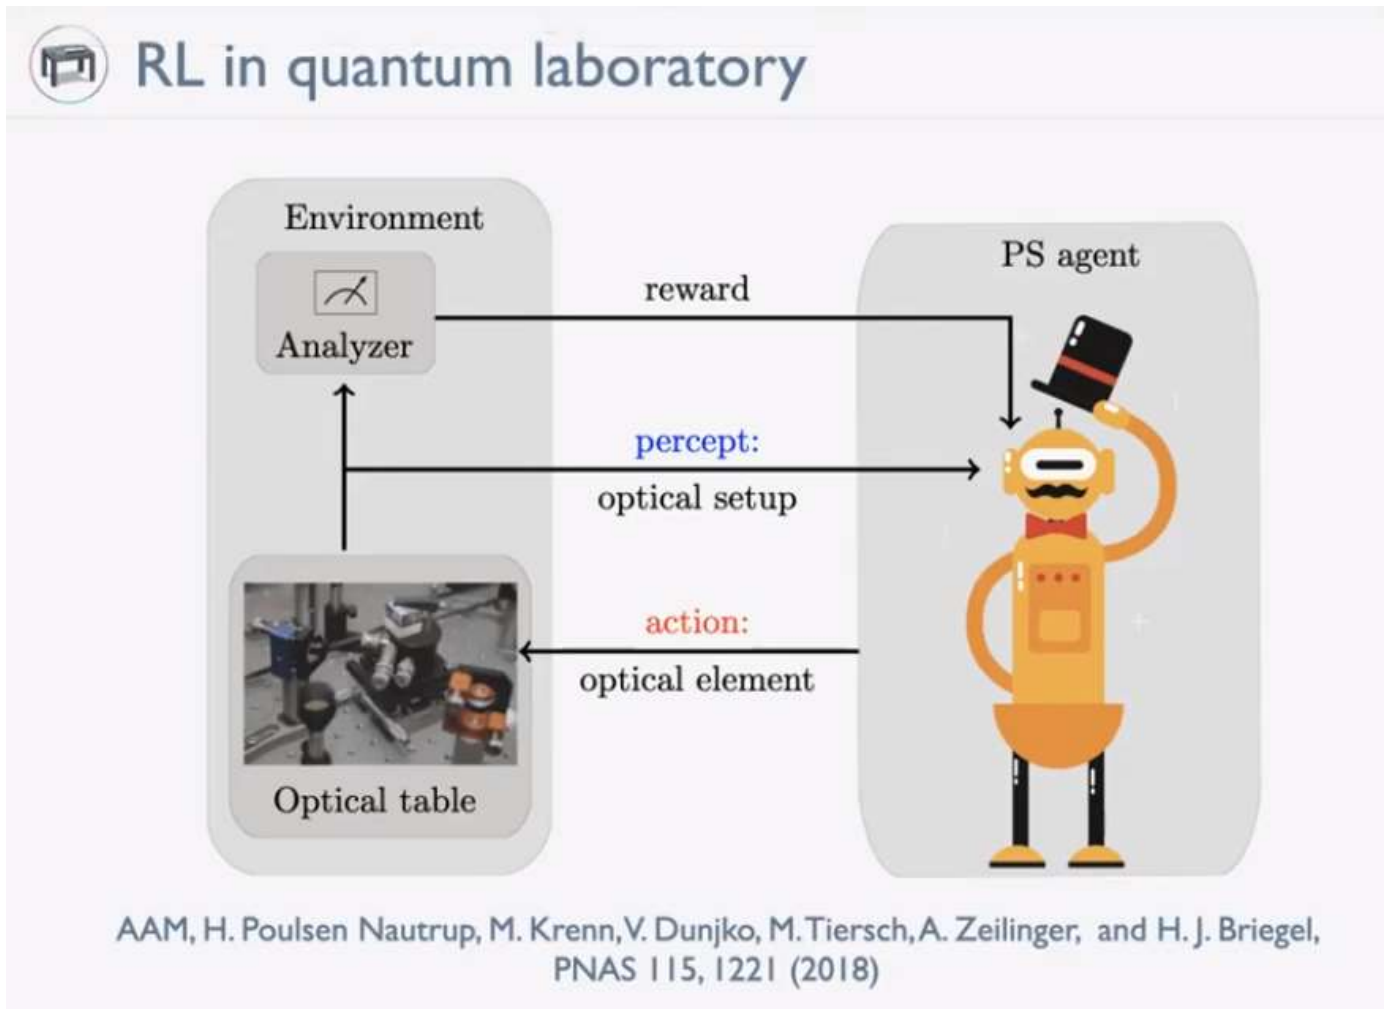
PS (cont.)

- Analytic solution for convergence of PS agents
 - a paper found
 - On the convergence of projective-simulation-based reinforcement learning in Markov decision processes
(<https://arxiv.org/pdf/1910.11914.pdf>).
- PS and its cost efficiency V.S. Q-learning and Sarsa
 - Easier to optimize
 - Fewer meta-param
- PS with meta-param controller
- PS with generalization clips

- PS and their applications



PS agent on Quantum System



- The task: creating photonic quantum experiment
 - actions: state transition
 - reward: given by analyzer (human)
- The model
 - naive + clip composition and deletion
- The experimental results
 - PS is efficient compared to previous methods
 - analysis of the memory clips (with strong connectivities)
 - rediscover useful subroutines!
 - conclusion (PS help to)
 - find interesting experiments
 - find shorter implementation
 - rediscover subroutines

Note for Day 4 (12/5)

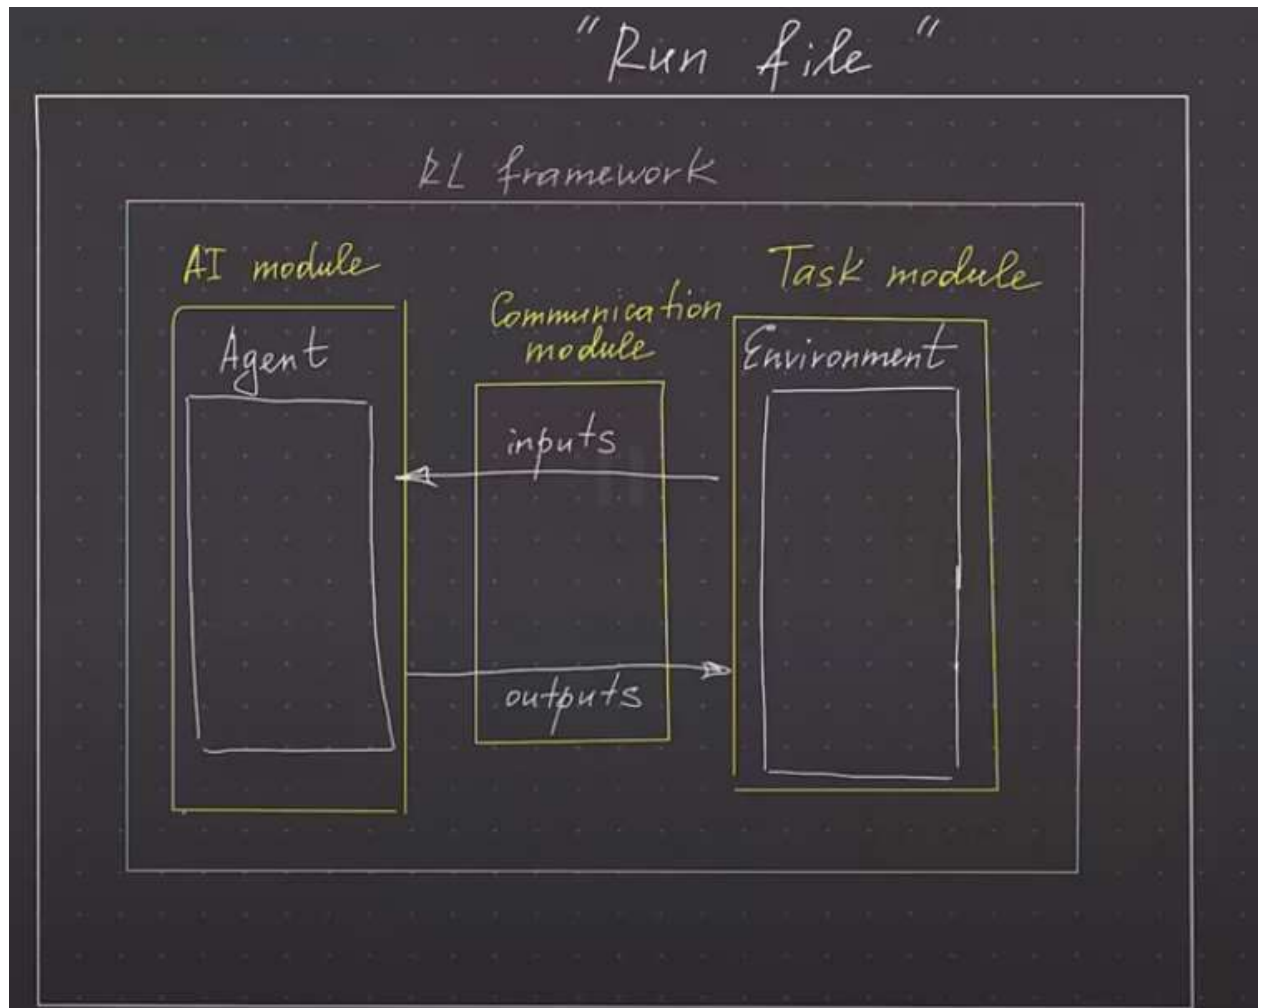
Overview

- PS code study

- PS applied to Quantum (practice)
- Quantum Enhanced PS

PS code study

- modularization



- 4 types of agents (categorized by clips)
 - basic (as in lecture)
 - flexible (states are unknown and can be added, indexed by hash function)
 - sparse (the sparse version of basic agent)
 - generalization (internal clips by state properties, specified by the user)
 - there is some interesting mathematical analysis, by is not covered
 - code is too much to go through line by line in the lecture
- Environments
 - grid world
 - open ai (API)
- Interaction module
 - interaction
- For more information
 - projection simulation with generalization

- benchmarking PS v.s. Q-learning/SARSA
- PS for experiment design
- ML for quantum communication
- RL for violating Bell's inequalities
- The code for the papers is the one that just covered with some twist and the papers are accepted immediately for the contributions to the community

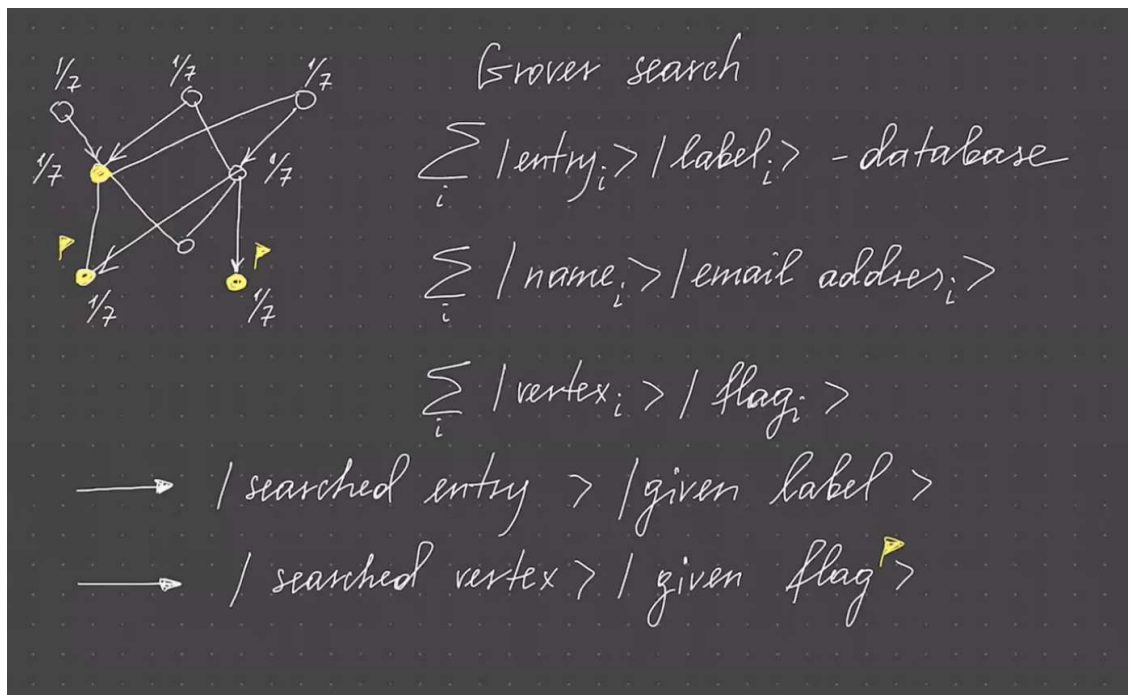
Note for Day 5 (12/11)

Overview

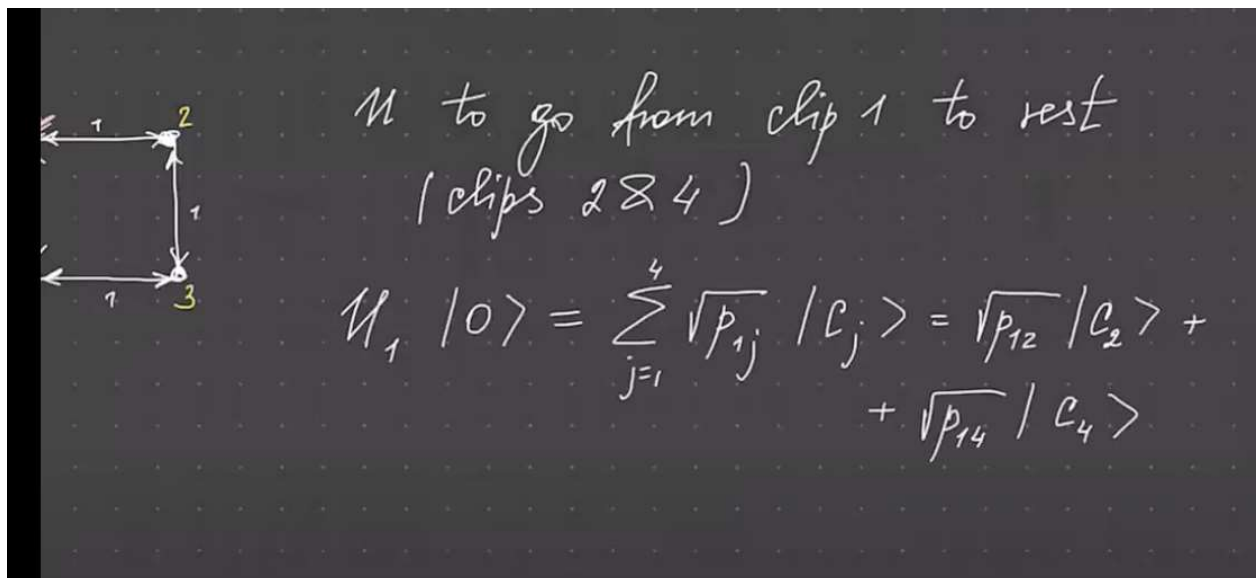
- Quantum Devices to Speed-up PS (Q for ML)
- Quantum Communication Tasks with PS (ML for Q)
 - Generalized Bell Test

Quantum Speed up for PS

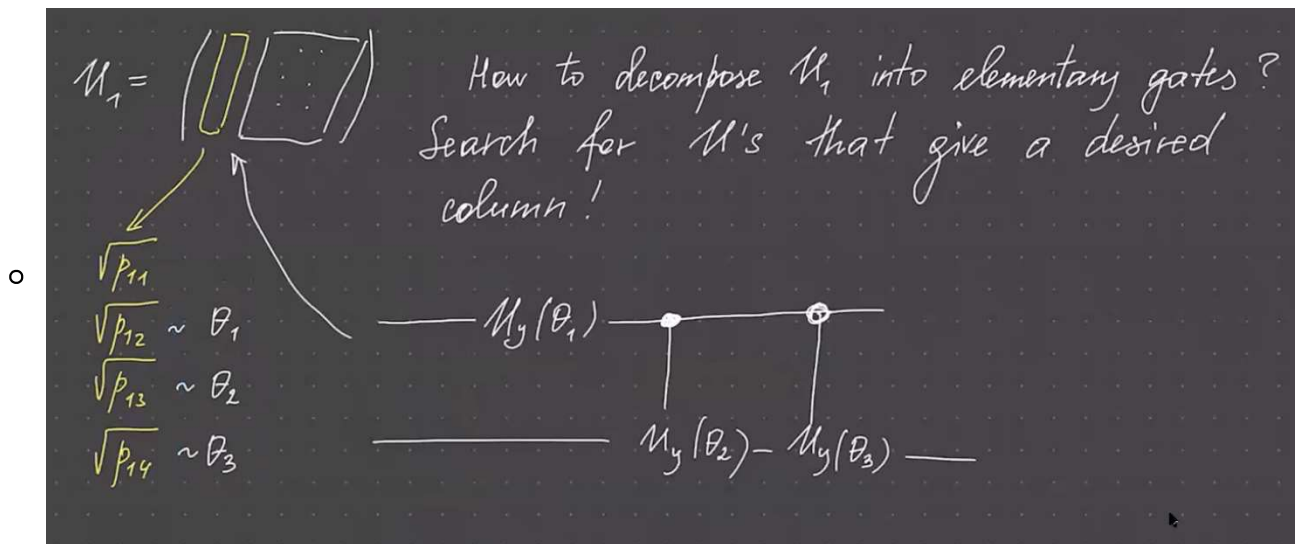
- Refs
 - Quantum Speedup for Active Learning Agents
 - an obvious application of Grover's Search
 - Quantum-enhanced deliberation of learning agents using trapped ions
 - Coherent controlization using superconducting qubits
- quadratic speed up for
 - it takes random walk $t = \frac{1}{\delta}$ to converge to the uniform distribution of all vertex
 - it takes quantum walk $t = \frac{1}{\sqrt{\delta}}$ to converge to uniform distribution of all vertex
 - if the distribution is a uniform distribution, with ϵ prob we will find the desired state
 - this can speed up using amplitude amplification too



- Question:
 - Quantum Walk first or Projective Simulation First? (quantum motivated PS?)
- Amplitude Amplification, in this case, will preserve the \hbar value's effect so that end with a high reward is more likely to be observed
- Quantum Walk Formalism (and an example)
 - $U_i|0\rangle = \sum_j \sqrt{p_{ij}}|j\rangle$



- Circuit Synthesis



- Q: how to implement control-U?
- I am not of physics background, this part is not clear for me. If this will be a crucial part of my future research, I will fill it up.
- Conclusion
 - Any decision making on N clips can be implemented on $\log(N)$ qubits
 - We can prepare stationary distribution ($Tp = p$) faster (w.r.t. spectral gap δ)
 - classical: $t = \frac{1}{\delta}$
 - quantum: $t = \frac{1}{\sqrt{\delta}}$
 - We can find desired actions faster (w.r.t. the fraction of desired state ϵ)
 - classical: $t = \frac{1}{\epsilon}$
 - quantum: $t = \frac{1}{\sqrt{\epsilon}}$

Quantum Communication Tasks with PS

- bell test

Learning to violate Bell inequality

State preparation

Measurement settings A_0, A_1

Alice

+1/-1

Measurement settings B_0, B_1

Bob

+1/-1

The experiment is repeated many times to estimate 16 probabilities:

$p(+1 + 1 | A_0 B_0)$

$p(+1 - 1 | A_0 B_0)$

$p(-1 + 1 | A_0 B_0)$

$p(-1 - 1 | A_1 B_1)$

...

CHSH value:

$$\beta = \sum_{x,y=0}^1 (-1)^{xy} (p(a = b | A_x B_y) - p(a \neq b | A_x B_y))$$

Any improvement in β is crucial for implementing DIQKD

Universität Basel Alexey Melnikov Machine learning for setting up quantum experiments 1

- generalized bell test

Getting the best CHSH value with optics

Lecture Day 5-2.key

A way to find the best CHSH value:

- Write down a general Bogolyubov transformation on n bosonic modes

$$G = V \circ \text{SMS}_g \circ U \circ D_\alpha$$

Alice

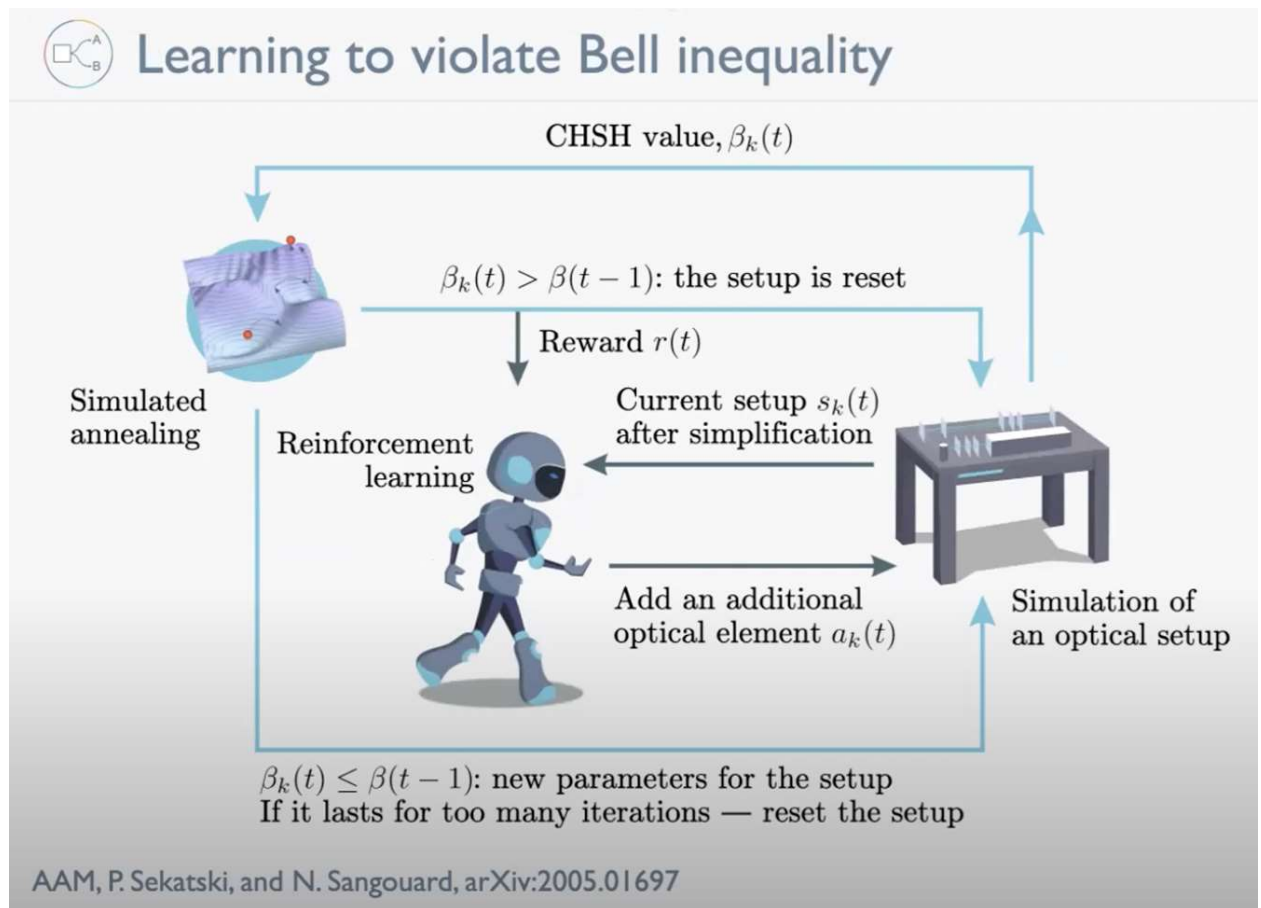
Bob

- Decompose linear optics part into beam splitters and phase shifters
- Optimize the parameters of all optical elements

AAM, P. Sekatski, and N. Sangouard, arXiv:2005.01697

- my observation: this different from previous examples, the space is continuous

- the framework



- simulated annealing
 - to find an extreme value in a n dimensional space
 - 1. random initialization
 - 2. choose by prob which to change
 - 3. choose by prob change how much
 - 4. get the new β'
 - $\beta' < \beta$, accept the change
 - $\beta' \leq \beta$, accept the change with prob $e^{\frac{\beta' - \beta}{\tau}}$

Note for Day 6 (12/12)

- For (continuous-time) quantum walk, I actually did a survey for a related paper, the link is shown below
 - https://github.com/EazyReal/QCQI2020fall/blob/main/QCQI_final_Oral_Version_.pdf
[f.\(https://github.com/EazyReal/QCQI2020fall/blob/main/QCQI_final_Oral_Version_.pdf\)](https://github.com/EazyReal/QCQI2020fall/blob/main/QCQI_final_Oral_Version_.pdf)
- For CNNs, I have taken DNN courses so that I am quite familiar with them
- The paper tries to predict when will a quantum speedup happen (graph \rightarrow speed-up?)

- o Predicting quantum advantage by quantum walk with convolutional neural networks

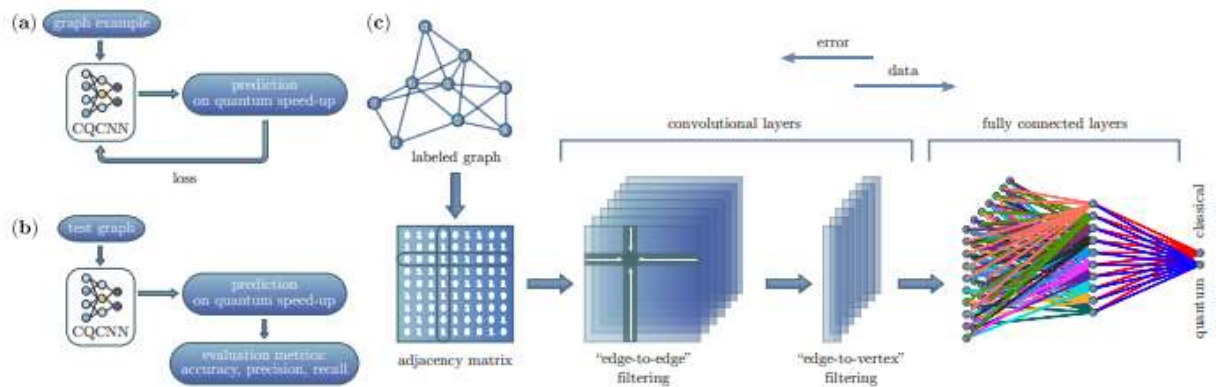


FIG. 2. The machine learning approach that is used for predicting the quantum speedup. (a) A process of training CQCNN. (b) A process of testing CQCNN. (c) A scheme of the CQCNN architecture. The neural network takes a labeled graph in form of an adjacency matrix as an input. This input is then processed by convolutional layers with graph-specific "edge-to-edge" and "edge-to-vertex" filters (see Methods). The convolutional layers are connected with fully-connected layers that finally classify the input graph. The number of layers is the same for all graph sizes. Data and error propagation are shown with arrows.

Materials

- Quantum enhancements for deep reinforcement learning in large spaces
- On the convergence of projective-simulation-based reinforcement learning in Markov decision processes
- Predicting quantum advantage by quantum walk with convolutional neural networks

HW2 Problem Set

- For the code, please visit [my GitHub repo \(https://github.com/EazyReal/Quantum-Machine-Learning-2020-fall\)](https://github.com/EazyReal/Quantum-Machine-Learning-2020-fall).
 - o Grid World
 - For finding the optimal $\eta = 0.1$, visit `run_grid.py`
 - For visualizations, visit `Single Agent in Grid World.ipynb`
 - o Grover's Diffusion Operator Synthesis
 - For the environment implementation, visit `environments/env_quantum_circuit_synthesis.py`
 - For finding the optimal $\eta = 0.9$, visit `run_grover.py`
 - For visualizations, visit `Single Agent doing Grover's Diffusion Operator.ipynb`

Q1

Go through the material which we considered during the lectures. Think of what is not clear and ask questions.

A1

- Quantum speedup of projection simulation with quantum walk is interesting, and here is my question. Do the authors use projection simulation to make RL a random walk problem so that we can have a quantum speedup? Or the authors come up with projective simulation first and find that they are random walks?
- Advice for the code (for `rl_framework.py`), you may try `**kwargs` syntax sugar in Python to pass the configuration as parameters!

Q2

- Run the PS code from Github on your computers. Understand how to set up the code to run the Grid World environment and solve it with PS. Try three implementations: flexible, basic, sparse.
- Send me in Direct Message the outputs you obtain in running a single PS agent in the Grid World environment.
 - Plot reward vs. trials curve.
 - Plot number of steps vs. trials curve.

A2

Test How to Run the Code and Solve GridWorld with the `base` agent

- Simply run `run.py` with a proper configuration and the results (performance for each eta average over agents, or something like this) will be stored in a `results` folder
 - `env_name` : which environment to use
 - `env_config` : the configuration of the environment
 - `multiple_agents` determines how the Env and Agent s interact (1-1 or 1-many)
 - `num_agents` : how many agents to take average on (in the `multiple_agents = False` scene), or how many agents are in the environment (in the `multiple_agents = True` scene)
 - `agent_name` : choose the type of agent you want to use
 - `max_num_trials` : how many trials do you want an agent to have
 - `max_steps_per_trial` : limit the number of steps in each trial (to avoid a trial that is too long or even an infinite loop)
- I improve the code with some additional features
 - learning curves for each eta
 - `run.py` and `rl_framework.py`
 - pass walls to GridWorld environment constructor with `env_config`
 - `run.py` and `env_grid_world.py`
 - apply `tqdm` to the loop to visualize the progress of the tasks


```
PS C:\Users\Maxwell Lin\Documents\projectivesimulation-master> python .\run_grover.py
0%|          | 0/11 [00:00<?, ?it/s]
10%|         | 1/10 [02:10<19:37, 130.83s/it]
```

- There's more that can be done if have time (But I am occupied with other things)
 - save the best agent of the best eta as an object, so that we can experiment with it

```
In [17]: performance = np.loadtxt('./results'+'/param_performance', delimiter=',')
```

```
In [18]: plt.plot(performance)
```

```
Out[18]: [matplotlib.lines.Line2D at 0x1b02b624610]
```

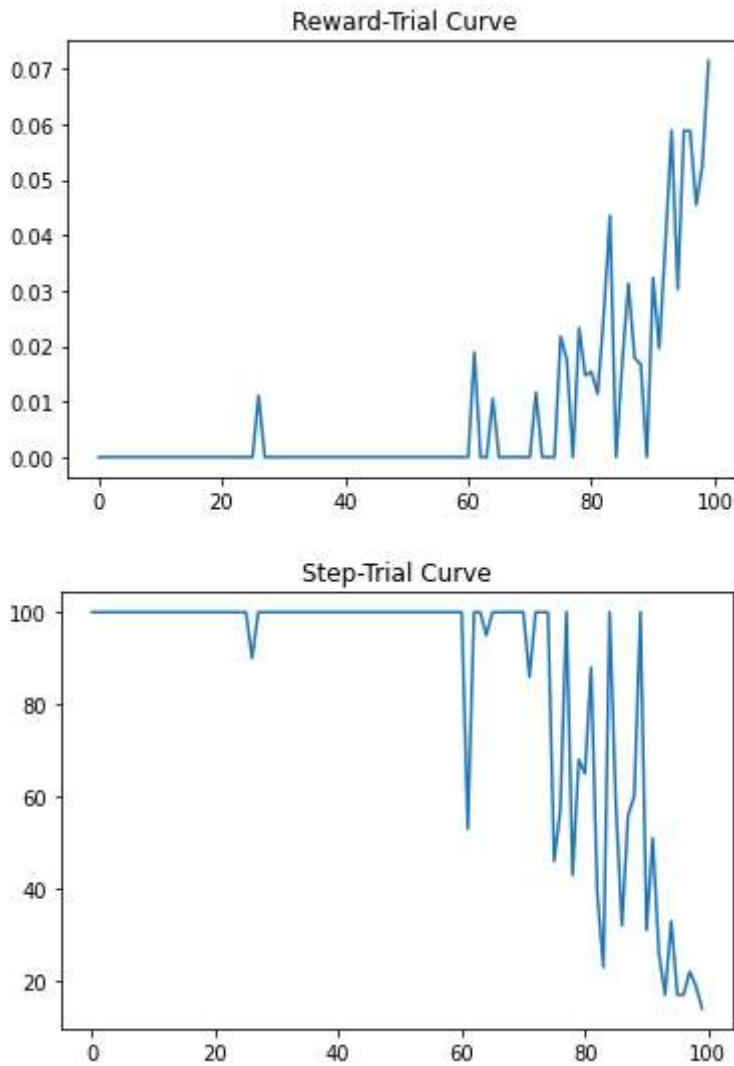


- By running the code for grid world environment (with my walls-passing code) and reading out the `average_param_performance` nd-array, we can see that the optimal η according to this experiment is $\in [0.1, 0.2]$
- The `sparse` (sparse matrix), `flexible` (dictionary of states) agents differ from `basic` agent only in the way they store states, and the states they stored should be identical, so I think there is no reason to try the same computational consuming process to find η s from the same distribution again. However, I did test the correctness of the codes with a smaller setting. I will omit this part for the quantum circuit experiment.

Results for the GridWorld problem with a single PS agent

- I modify the `simple_interaction.py` and plot the resulting data
 - the code can be found in the `Single Agent in Grid World.ipynb` in my GitHub repo
- The configuration is as follow
 - agent: `basic` PS agent with softmax and $\eta = 0.1$ (delayed reward), w/o reflections or gamma damping (forgetting)

- environment: the figure provided in the slide



Q3

Program a toy task environment. And solve it with PS (try basic, flexible, and sparse).

- Actions Implement quantum gates on a register of two qubits.
 - action #1: Hadamard gate on qubit 1
 - action #2: Hadamard gate on qubit 2
 - action #3: X gate on qubit 1
 - action #4: X gate on qubit 2
 - action #5: Z gate on qubit 1
 - action #6: Z gate on qubit 2
 - action #7: CNOT gate on qubits 1-2
- Environmental states
 - Current sequence of actions within current trial.
- Reward

- Reward of +1 is given for arriving to an environmental state that corresponds to a quantum circuit implementing Grover diffusion operator.
- In this task provide:
 - reward vs. trials curve
 - quantum circuit length vs. trials curve
 - draw circuits to which agent converges

A3

Implementation of the environment


There are two ways I come up with that can implement the environment

- direct calculation with numpy
 - $\text{np.kron}(A, B) = A \otimes B$
 - $A @ B = AB$
- "do not redo" with the existence of qiskit package
 - Unitary Simulator provided in Qiskit Aer
[.https://qiskit.org/documentation/stubs/qiskit.providers.aer.UnitarySimulator.html#qiskit.providers.aer.UnitarySimulator](https://qiskit.org/documentation/stubs/qiskit.providers.aer.UnitarySimulator.html#qiskit.providers.aer.UnitarySimulator)

```
backend_sim = Aer.get_backend('unitary_simulator')
job_sim = execute([qc1, qc2], backend_sim)
result_sim = job_sim.result()
unitary1 = result_sim.get_unitary(qc1)
unitary2 = result_sim.get_unitary(qc2)
np.allclose(unitary1, unitary2)
```

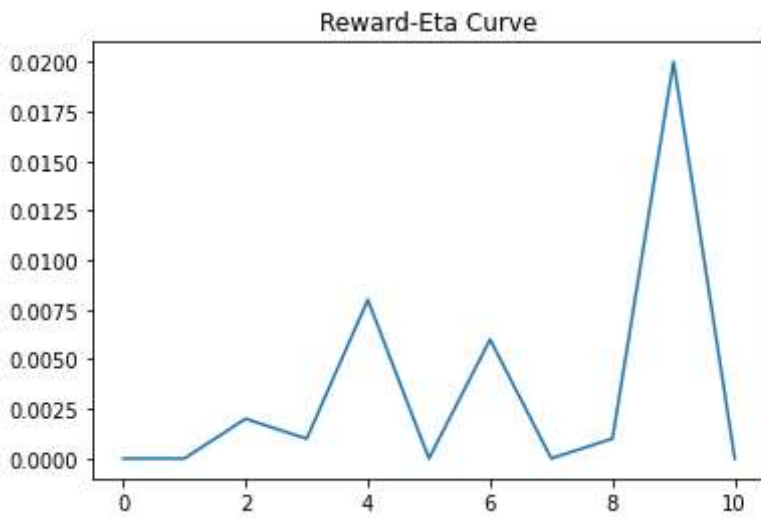
I chose the second one for simplicity, the "do not redo" principle and the convenience to draw a circuit. The code can be found in my Github repo.

The target circuit is represented by a common implementation below.

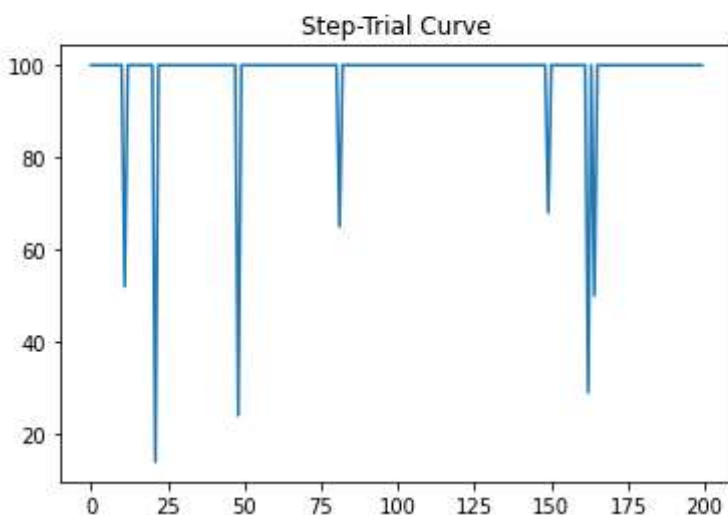



Solving for Optimal η with the `flexible` agent and plot the learning/step-trial curve

To solve for optimal η , I will use the `flexible` agent for the infinite possibilities of quantum circuits without giving a depth limit. Alternatively, one can try using a `sparse` agent with a pre-specified limit of depth.



By running the modified `run.py` and reading out the `average_param_performance` nd-array, we can see that the optimal η according to this experiment is about 0.9.



However, I do not observe an obvious "learning" (or say generalization) behavior of the agents since there is no trend of the agents becoming better over trials. I also tried many times and tried with different eta and the results are similar. I think this is because that the state ("current circuit composition") is too sparse and the agents fail to recognize some


"circuit identities" (e.g. $XX = I$, $ZZ = I$, $HH = I$, $HZH = X$, etc.), so unlike in the grid world problem, the agents are unable to tell if a state is equivalent to a state they visited before in terms of unitary.

There are three ways to improve this


- treat "equivalence class of unitaries" as states directly
 - this violate the spec of the homework
- increase the number of reflections.
 - no direct support for number of reflections in the `flexible` agent
- increase the number of trials.
 - I tried number of trials = 1000, does not help

Draw circuits to which agents converge


In [38]: `saved_circuit.draw()`




In [16]: `saved_circuit.draw()`



In [56]: `saved_circuit.draw()`



In [29]: `saved_circuit.draw()`



- Despite that learning did not happen, the agent did get some circuits that is equivalent to the Grover's diffusion operator.
- One can verify the circuits are equivalent to the Grover's Diffusion Operator
 - with the circuit identities: $HH = I$, $XX = I$, $ZZ = I$, $HZH = X$